

# Gibbs Sampling for a mixture of normals

Matthew Stephens

2016-06-01

## Pre-requisites

You should know about Gibbs sampling and mixture models, and be familiar with Bayesian inference for the normal mean and for the two class problem.

## Overview

We consider using Gibbs sampling to perform inference for a normal mixture model,

$$X_1, \dots, X_n \sim f(\cdot)$$

where

$$f(\cdot) = \sum_{k=1}^K \pi_k N(\cdot; \mu_k, 1).$$

Here  $\pi_1, \dots, \pi_K$  are non-negative and sum to 1, and  $N(\cdot; \mu, \sigma^2)$  denotes the density of the  $N(\mu, \sigma^2)$  distribution.

Recall the latent variable representation of this model:

$$\Pr(Z_j = k) = \pi_k$$

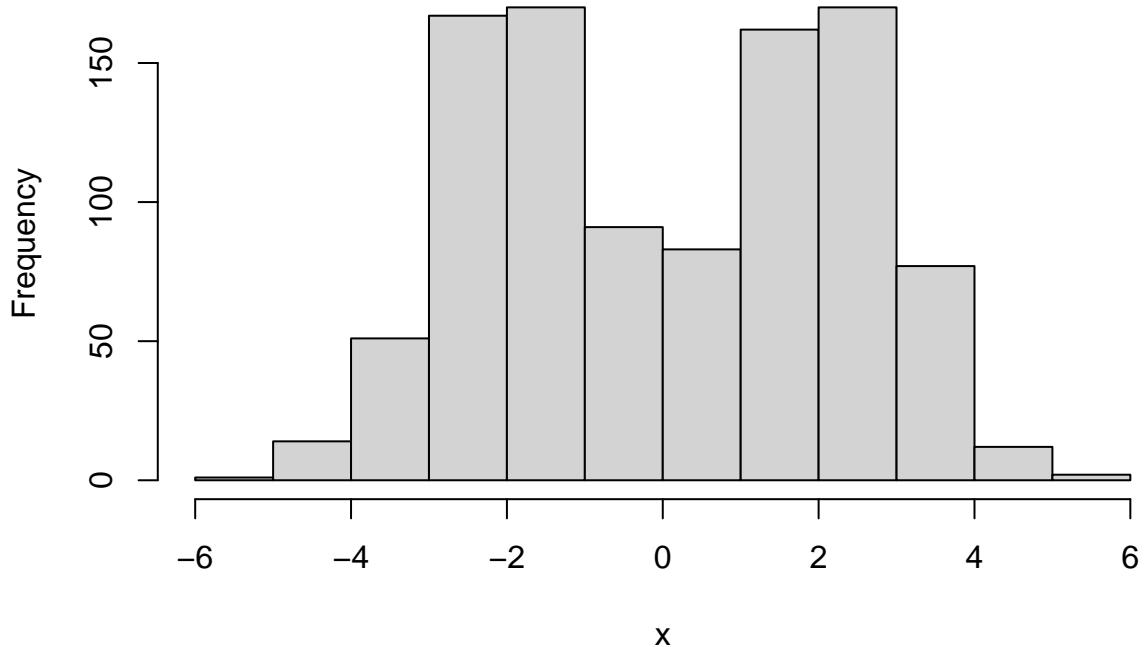
$$X_j | Z_j = k \sim N(\mu_k, 1)$$

To illustrate, let's simulate data from this model:

```
set.seed(33)

# generate from mixture of normals
#' @param n number of samples
#' @param pi mixture proportions
#' @param mu mixture means
#' @param s mixture standard deviations
rmix = function(n,pi,mu,s){
  z = sample(1:length(pi), prob=pi, size=n, replace=TRUE)
  x = rnorm(n, mu[z], s[z])
  return(x)
}
x = rmix(n=1000, pi=c(0.5,0.5), mu=c(-2,2), s=c(1,1))
hist(x)
```

## Histogram of x



## Gibbs sampler

Suppose we want to inference for the parameters  $\mu, \pi$ . That is, we want to sample from  $p(\mu, \pi|x)$ . We can use a Gibbs sampler. However, to do this we have to augment the space to sample from  $p(z, \mu, \pi|x)$ , not only  $p(\mu, \pi|x)$ .

Here is the algorithm in outline:

- sample  $\mu$  from  $\mu|x, z, \pi$
- sample  $\pi$  from  $\pi|x, z, \mu$
- sample  $z$  from  $z|x, \pi, \mu$

The point here is that all of these conditionals are easy to sample from.

## Code

```

normalize = function(x){return(x/sum(x))}

#' @param x an n vector of data
#' @param pi a k vector
#' @param mu a k vector
sample_z = function(x,pi,mu){
  dmat = outer(mu,x, "-") # k by n matrix, d_kj =(mu_k - x_j)
  p.z.given.x = as.vector(pi) * dnorm(dmat,0,1)
  p.z.given.x = apply(p.z.given.x,2,normalize) # normalize columns
  z = rep(0, length(x))
  for(i in 1:length(z)){
    z[i] = sample(1:length(pi), size=1, prob=p.z.given.x[,i], replace=TRUE)
  }
}

```

```

    return(z)
}

#' @param z an n vector of cluster allocations (1...k)
#' @param k the number of clusters
sample_pi = function(z,k){
  counts = colSums(outer(z,1:k,FUN=="=="))
  pi = gtools::rdirichlet(1,counts+1)
  return(pi)
}

#' @param x an n vector of data
#' @param z an n vector of cluster allocations
#' @param k the number o clusters
#' @param prior.mean the prior mean for mu
#' @param prior.prec the prior precision for mu
sample_mu = function(x, z, k, prior){
  df = data.frame(x=x,z=z)
  mu = rep(0,k)
  for(i in 1:k){
    sample.size = sum(z==i)
    sample.mean = ifelse(sample.size==0,0,mean(x[z==i]))

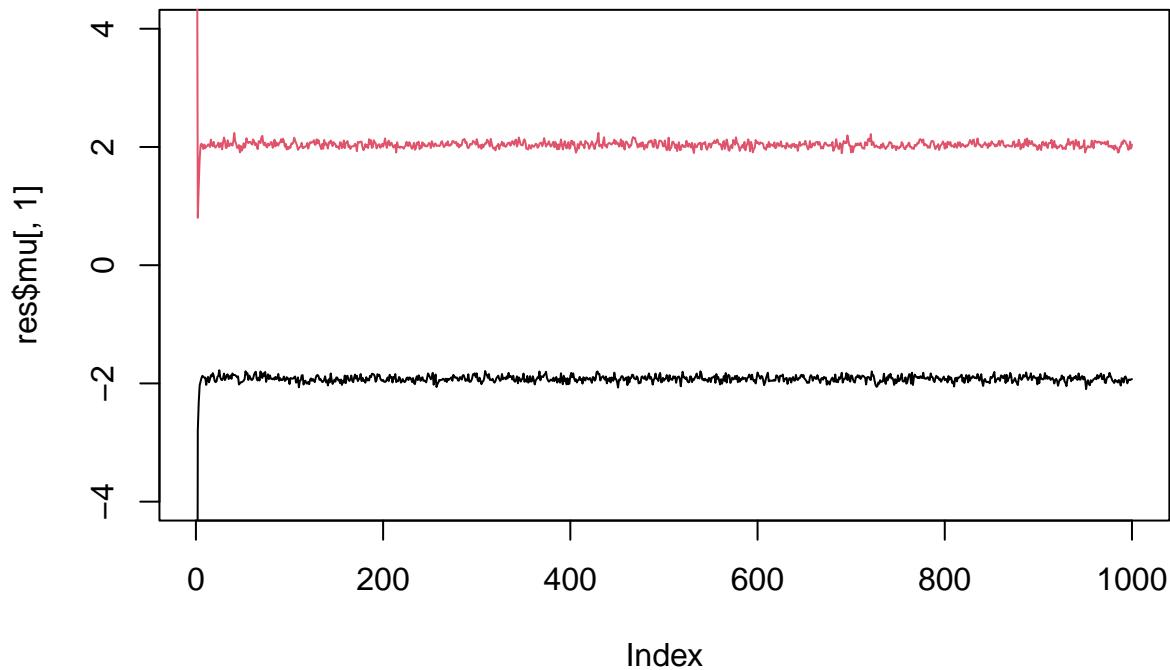
    post.prec = sample.size+prior$prec
    post.mean = (prior$mean * prior$prec + sample.mean * sample.size)/post.prec
    mu[i] = rnorm(1,post.mean,sqrt(1/post.prec))
  }
  return(mu)
}

gibbs = function(x,k,niter =1000,muprior = list(mean=0,prec=0.1)){
  pi = rep(1/k,k) # initialize
  mu = rnorm(k,0,10)
  z = sample_z(x,pi,mu)
  res = list(mu=matrix(nrow=niter, ncol=k), pi = matrix(nrow=niter,ncol=k), z = matrix(nrow=niter, ncol=k))
  res$mu[1,]=mu
  res$pi[1,]=pi
  res$z[1,]=z
  for(i in 2:niter){
    pi = sample_pi(z,k)
    mu = sample_mu(x,z,k,muprior)
    z = sample_z(x,pi,mu)
    res$mu[i,] = mu
    res$pi[i,] = pi
    res$z[i,] = z
  }
  return(res)
}

```

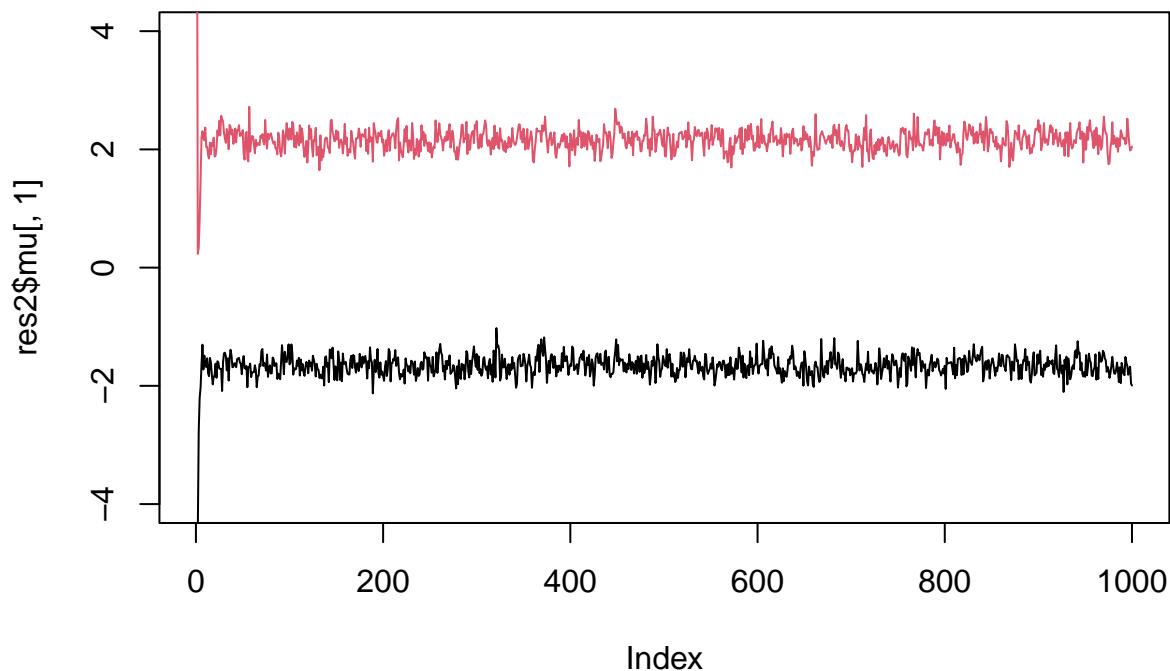
Try the Gibbs sampler on the data simulated above. We see it quickly moves to a part of the space where the mean parameters are near their true values (-2,2).

```
res = gibbs(x, 2)
plot(res$mu[, 1], ylim=c(-4, 4), type="l")
lines(res$mu[, 2], col=2)
```



If we simulate data with fewer observations we should see more uncertainty

```
x = rmix(100, c(0.5, 0.5), c(-2, 2), c(1, 1))
res2 = gibbs(x, 2)
plot(res2$mu[, 1], ylim=c(-4, 4), type="l")
lines(res2$mu[, 2], col=2)
```

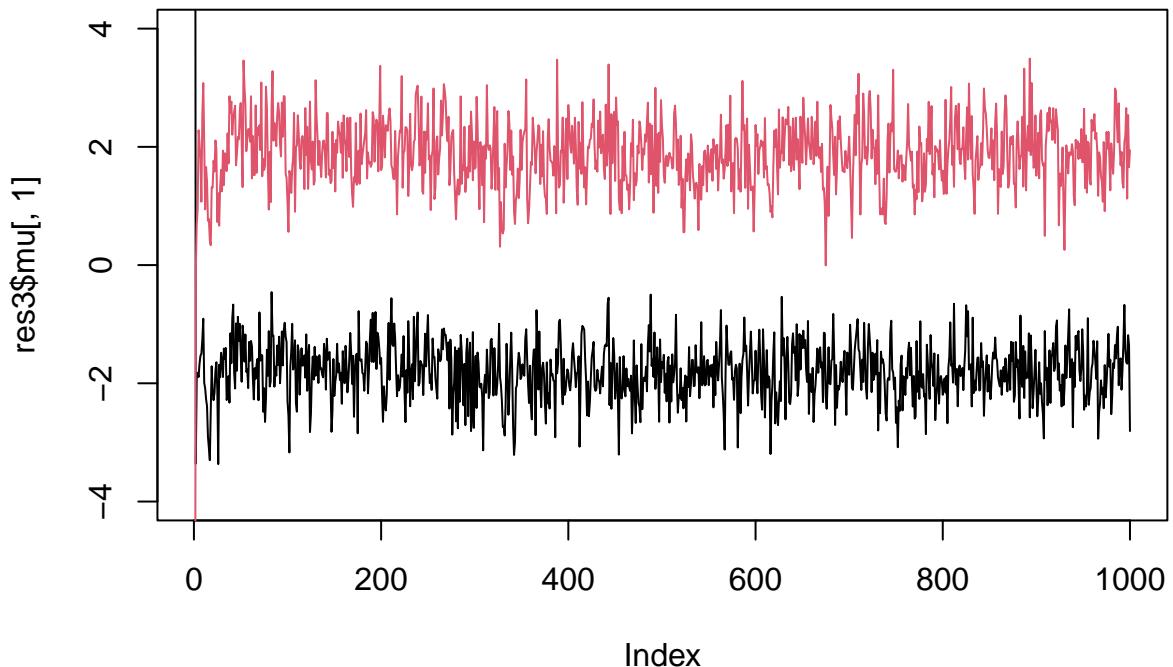


And fewer observations still...

```

x = rmix(10,c(0.5,0.5),c(-2,2),c(1,1))
res3 = gibbs(x,2)
plot(res3$mu[,1],ylim=c(-4,4),type="l")
lines(res3$mu[,2],col=2)

```



And we can get credible intervals (CI) from these samples (discard the first few samples as “burn-in”).

For example, to get 90% posterior CIs for the mean parameters:

```

quantile(res3$mu[-(1:10),1],c(0.05,0.95))

##      5%      95%
## -2.645 -1.004

quantile(res3$mu[-(1:10),2],c(0.05,0.95))

##      5%      95%
## 0.940  2.777

```