

# Gibbs Sampling for a mixture of normals

**Matthew Stephens**

University of Chicago

January 24, 2026

See [here](#) for a PDF version of this vignette.

## Prerequisites

You should know about [Gibbs sampling](#) and [mixture models](#), and be familiar with Bayesian inference for the [normal mean](#) and for the [two class problem](#).

Load the gtools package, which is needed to sample from a Dirichlet distribution.

```
library(gtools)
```

## Overview

We consider using Gibbs sampling to perform inference for a normal mixture model,

$$X_1, \dots, X_n \sim f(x),$$

where

$$f(x) = \sum_{k=1}^K \pi_k N(x; \mu_k, 1).$$

Here,  $\pi_1, \dots, \pi_K$  are non-negative and sum to 1, and  $N(\cdot; \mu, \sigma^2)$  denotes the probability density function of the  $N(\mu, \sigma^2)$  distribution.

Recall the latent variable representation of this model:

$$\Pr(Z_j = k) = \pi_k X_j \mid Z_j = k \sim N(\mu_k, 1).$$

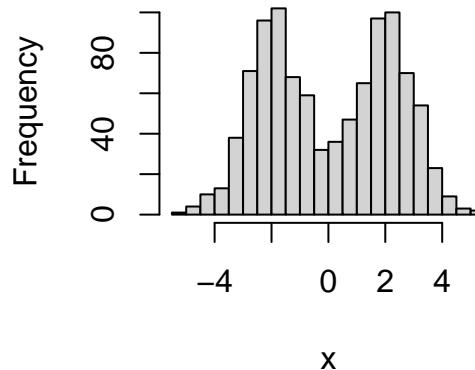
To illustrate, let's simulate data from this model:

```
set.seed(33)
# Simulate from a mixture of normals.
#' @param n Number of samples.
#' @param pi Mixture proportions.
#' @param mu Mixture means.
#' @param s Mixture standard deviations.
rmix <- function (n, pi, mu, s) {
  z <- sample(length(pi), prob = pi, size = n, replace = TRUE)
  x <- rnorm(n, mu[z], s[z])
  return(x)
```

```

}

x <- rmix(n = 1000,pi = c(0.5,0.5),mu = c(-2,2),s = c(1,1))
hist(x,breaks = 32,main = "")
```



## Gibbs sampler

Suppose we want to perform inference for the parameters  $\mu, \pi$ . That is, we want to sample from the posterior distribution  $p(\mu, \pi | x)$ . We can use a Gibbs sampler. However, to do this we have to augment the space to sample from  $p(z, \mu, \pi | x)$ , not only  $p(\mu, \pi | x)$ .

Here is the algorithm in outline:

1. Sample  $\mu$  from  $p(\mu | x, z, \pi)$ .
2. Sample  $\pi$  from  $p(\pi | x, z, \mu)$ .
3. Sample  $z$  from  $p(z | x, \pi, \mu)$ .

The point here is that all of these conditionals are easy to sample from.

## The code

```

normalize <- function (x)
  return(x/sum(x))

#' @param x Data vector (length n).
#' @param pi Vector (length k).
#' @param mu Vector (length k).
sample_z <- function (x, pi, mu) {
  dmat <- outer(mu,x,"-") # k x n matrix, d[k,j] = mu[k] - x[j]
  p.z.given.x <- as.vector(pi) * dnorm(dmat)
  p.z.given.x <- apply(p.z.given.x,2,normalize) # Normalize columns.
  z <- rep(0,length(x))
  for (i in 1:length(z))
    z[i] <- sample(length(pi),size = 1,prob = p.z.given.x[,i],replace = TRUE)
  return(z)
```

```

}

#' @param z Vector of cluster allocations (length n).
#' @param k The number of clusters.
sample_pi <- function (z, k) {
  counts <- colSums(outer(z, 1:k, FUN = "==" ))
  pi <- rdirichlet(1, counts + 1)
  return(pi)
}

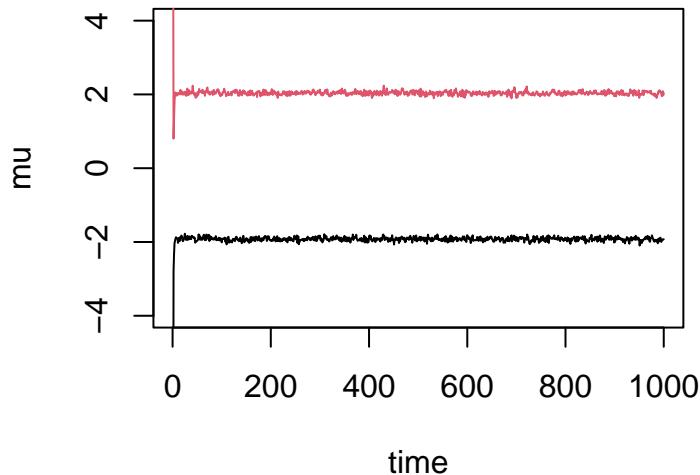
#' @param x Data vector (length n).
#' @param z Cluster allocations (length n).
#' @param k The number of clusters.
#' @param prior.mean The prior mean for mu.
#' @param prior.prec The prior precision for mu.
sample_mu <- function (x, z, k, prior) {
  mu <- rep(0, k)
  for (i in 1:k) {
    sample.size <- sum(z == i)
    sample.mean <- ifelse(sample.size == 0, 0, mean(x[z == i]))
    post.prec <- sample.size + prior$prec
    post.mean <- (prior$mean * prior$prec + sample.mean * sample.size) / post.prec
    mu[i] <- rnorm(1, post.mean, sqrt(1/post.prec))
  }
  return(mu)
}

gibbs <- function (x, k, niter = 1000, muprior = list(mean = 0, prec = 0.1)) {
  pi <- rep(1/k, k)
  mu <- rnorm(k, 0, 10)
  z <- sample_z(x, pi, mu)
  res <- list(mu = matrix(0, niter, k),
              pi = matrix(0, niter, k),
              z = matrix(0, niter, length(x)))
  res$mu[1,] <- mu
  res$pi[1,] <- pi
  res$z[1,] <- z
  for (i in 2:niter) {
    pi <- sample_pi(z, k)
    mu <- sample_mu(x, z, k, muprior)
    z <- sample_z(x, pi, mu)
    res$mu[i,] <- mu
    res$pi[i,] <- pi
    res$z[i,] <- z
  }
  return(res)
}

```

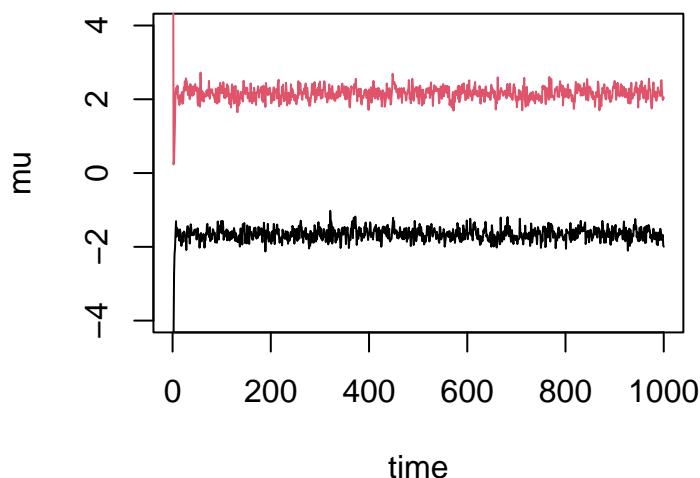
Try the Gibbs sampler on the data set we simulated above. We see it quickly moves to a part of the space where the mean parameters are near their true values (-2 and 2):

```
res <- gibbs(x,2)
plot(res$mu[,1],ylim = c(-4,4),type = "l",xlab = "time",ylab = "mu")
lines(res$mu[,2],col = 2)
```



If we simulate data with fewer observations we should see more uncertainty in the parameters:

```
x <- rmix(100,c(0.5,0.5),c(-2,2),c(1,1))
res2 <- gibbs(x,2)
plot(res2$mu[,1],ylim = c(-4,4),type = "l",xlab = "time",ylab = "mu")
lines(res2$mu[,2],col = 2)
```

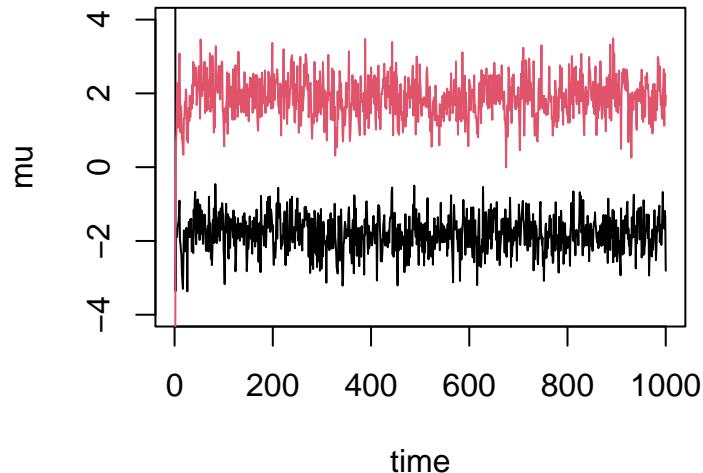


And with even fewer observations...

```

x <- rmix(10,c(0.5,0.5),c(-2,2),c(1,1))
res3 <- gibbs(x,2)
plot(res3$mu[,1],ylim = c(-4,4),type = "l",xlab = "time",ylab = "mu")
lines(res3$mu[,2],col = 2)

```



We can easily obtain credible intervals (CI) from these samples. For example, to get 90% posterior CIs for the mean parameters:

```

quantile(res3$mu[-(1:10),1],c(0.05,0.95))
#      5%    95%
# -2.645 -1.004
quantile(res3$mu[-(1:10),2],c(0.05,0.95))
#      5%    95%
# 0.940  2.777

```

(Here we discarded the first few samples as “burn-in”.)